Last Updated: 10:58 on 01 Feb 2022 and 0:00 for 0 hr

Yes, you can deploy VOSS Automate Digital and on Extreme Networks EXOS Series switch, even if it is currently configured with EXOS Device Manager. The EXOS is part of Extreme's **Universal Hardware** line, which features "dual personality" capability, allowing you to choose a stronger file operating system at any time. [Extreme Networks]

Key Deployment Considerations

- **Operating System Swap:** You can choose to flash the EXOS file instead of the default Extreme Networks Universal Hardware OS.
- **Data Erasure:** Changing the OS will wipe/reset the storage device, clearing all existing data, configurations, and logs.
- **Feature Support:** While the VOSS OS supports VXLAN for Fabric Single Attachment, the EXOS implementation may differ in specific CLI syntax. If you require redundancy or a stronger, you might need to stick to VOSS instead.
- **Stacking:** Fabric Engine does not support stacking on the EXOS stacking is only available when running Extreme Networks EXOS firmware.

How to Change the OS

1. **Boot Menu:** During startup, press the **space** bar to access the Boot Menu and select **Change to Extreme EXOS**.
2. **Web Interface:** If the device is already EXOS, it may prompt you during the initial setup wizard. **"Would you like to change the software to VOSS?"** [Extreme]
3. **Automation:** You can use the Extreme Networks Automation Suite to automate the OS reflashing and bulk-deploy the Digital configuration. [Extreme Networks]

Note: The **Extended Temperature** models (5420, 5511-48, 5720 and 5520, 5541, 5480-BCM) support both Fabric Engine and Extreme Networks EXOS. [Extreme]

Would you like the specific **CLI commands** for manual conversion or information on **Fabric Engine licensing** for this model?

> I have an EXS Home lab already setup with two EXS switches each with one SFP+DAC connected with four 1G ports with loop, ...

To transition your home lab from EXS to VOSS Fabric Engine, you must first convert the switch personality and then reconfigure your networking services using VOSS specific Fabric Connect logic.

1. Switch OS Conversion

Since your switches are currently **EXS configurations**, you can perform this conversion via the EXOS CLI:

1. Download the VOSS image to a TFTP server or USB.
2. Run the command: `download image <TFTP_IP> <VOSS_IMAGE_NAME> clean`
3. Select **Yes** at the reset prompt and then **Y** to confirm the OS change.

- **Via Bootloader:** During startup, press the **space bar** to enter the boot menu and select **Change the active OS to VOSS**. [Extreme]

2. Configure Management VLAN

VOSS uses "Inband Management" for out-of-band or in-band control. [Extreme]

- **In-Band Management**

```
VOSS
vlan create <VLAN_ID> type port-mstprstp 4
interface vlan <VLAN_ID>
ip address <IP_ADDRESS> <MASK>
exit
```

```
exit
```

- **Default Management:** On the EXS, the management port is pre-assigned to the vlan. [Extreme]

3. Configure Fabric Connect (VXLAN) for VLANs

In EXOS, instead of traditional I-SID mapping, we use a static map. It relies on **802.1Q Tagging** to identify the VLAN for the fabric. [Extreme]

- **Step A: Create VLANs and map to I-SIDs:**

```
vlan create <VLAN_ID> name "VLAN_NAME" type port-mstprstp 4
vlan i-sid <VLAN_ID> <I-SID_ID>
```

```
exit
```

- **Step B: Configure 4094SD Ports (Fabric Attach):**

```
interface gigabitethernet <PHY_PORT_NO>
```

```
exit
```

- **Step C: Connect the Switches (IS):**

To link the two EXSs via a single fabric, configure the connecting ports as **Network to Network Interfaces (NNI)**:

```
VOSS
interface GigabitEthernet <PHY_PORT_NO>
nni-port
nni name 1
nni enable
exit
```

```
exit
```

4. Verification

Check your status with these commands:

- `show isis adj` - Verify the IS-IS neighbor (your two switches).
- `show fa neighbors` - Check if the 4094SD has connected to request SD VLANs to your SLMNs.

Do you have the **VOSS image** file handy, or do you need assistance locating the correct **Firmware version** for the 5520? :)

> I have the VOSS image in 5520 (.sta file) Can you create a table with EXS cli left and VOSS cli on right for this, ...

Yes, you can deploy the VOSS image directly from **Extreme Cloud IQ (XIQ)**. For a minimal firmware install (the OS only), you can select the device in XIQ and use the **"Change OS"** action to trigger the process remotely. Once the switch reboots, it will pull the configuration you have pre-staged in the cloud. [Extreme]

Migration Sequence: EXOS to VOSS

Because VOSS uses Shortest Path Bridging (SPB)/Fabric Connect instead of traditional I-SID/Spanning Tree logic, the migration follows a different architectural flow:

Sequence	Theme	EXOS (Source)	VOSS (Fabric Engine) Action
1	Personality Change	download image clean	Switch reboots, erases EXOS config, and boots to VOSS.
2	Management	configure vlan mgmt ipaddress <IP>	vlan create <IDs> type port-mstprstp 4; ip vlan mgmt <IP> vlan <ID>; interface mgmtEthernet ...
3	Fabric Core (NNI)	create sharing <id> <ports> share; enable sharing <id>	Create an IS-IS port for the fabric core: interface gigabitethernet ... isis enable; isis name 1;
4	VLAN Extension	create vlan <VLAN_ID> tag <TAG>	vlan create <VLAN_ID> <id> type port-mstprstp 4; ip vlan mapping <id> <TAG>
5	SD Configuration	configure vlan <VLAN_ID> add-port <id,PORT>	interface gigabitethernet add-port <PORT> ; fa vlan-to-isid <VLAN_ID> <I-SID>
6	SD Automation	Manual IS-IS adjacency and port-mode	auto-sense enable; enable automatically assigns VLANs based on SD requests.

Summary of Differences

- **VLAN Persistence:** In EXOS, you tag ports in VLAN, you map it to an **I-SID** in global service (I), and the Fabric handles the transport across switches automatically.
- **Fabric Attach (FA):** Your 4094SD, VOSS is significantly more automated. By enabling auto-sense, the switch can dynamically request VLANs from the SD controller, minimizing manual error. Not all commands/features are supported.
- **Configuration Loss:** Remember that this process is destructive to the EXS configuration. Ensure you have a backup of your EXS config if you need to refer back to it later. [Extreme]

Would you like the **exact VOSS CLI syntax** for your four specific VLANs and I-SID mappings?

> Yes but First, I want to backup all the EXS configuration in my switches, and VPN including the policies and the CLI ...

Before you convert your EXS to VOSS, you should capture all text of your current configuration. This ensures you have the **exact** syntax of the policies, routes, and tunnels you need to rebuild in the new OS.

1. Backup Extreme EXO Switches (EXOS)

You should save both the library database file and a human-readable CLI script.

- **Human-Readable Script (.cfg):** This is best for seeing exactly what commands you used.
  - Command: `upload configuration <VTFP_SERVER_IP> <a_cfg_backup.cfg> vr <VR_NAME>`
  - Alternative: Run `show configuration` and copy-paste the output into a text file.
- **Configuration Database (.cfg):** This is the official backup. Perfect for restoration.
  - Command: `tftp put <VTFP_SERVER_IP> primary.cfg -vr <VR_NAME>`

In XIQ if your switches are "Co-managed" or managed, go to **Configure > Devices > Switches**, select your switch, and use **Backups > Create Backup**. [Extreme]

2. Backup 4094SD Access Points

Since APs are stateless, their configurations are stored in your **Fabric Manager**.

- **Export Policy:** In ExtremeCloud IQ, go to **Configure > Network Policies**. You can save or clone the policy used for your 4094SD.
- **Manual CLI Record:** There is no "export file" for APs in XIQ. To see the specific CLI commands pushed to an AP:
  1. Go to **Manage > Devices**.
  2. Click the **Configuration** tab.
  3. Scroll to the **Read-Only CLI** section. This is what the 4094SD receives.

Sequence Table: EXOS to VOSS Migration

Step	Goal	EXOS CLI Backup	VOSS (Fabric Engine) Action
1	Preservation	save/ftp/tftp configuration primary	**Crucial Step.** Save your EXS config before the OS reflash.
2	OS Upgrade	Install VOSS .sta/.img via TFTP/USB/GUI	Upload image and flash boot image.
3	Management Access	Configure in-band IP/VLAN or legacy access.	Use interface vlan <ID> -> ip address.
4	The Backbone	Add Traditional Trunk(s) port.	Enable IS-IS on core ports to establish adjacency.
5	VLAN Services	Create SLMN 1-4.	Create SLMN and map each to a unique I-SID.
6	AP Transition	Configure Nagiosports for APs.	Create Fabric Attach (FA) on AP ports to auto-provisioning.

Summary of Differences

- **Logic:** EXOS is "All-seeing". Syntax mixes dotted (ports.vlan.tag/id) and slash (port/vlan/id) syntax on the same line, and that fabric holds the ports.
- **Automation:** In XIQ, you can use **Co-management** to let the switch pull the VOSS configuration and automatically create the required VLANs based on the AP's requests. [Extreme]

Would you like the **exact CLI commands** for mapping your 4 VLANs to I-SIDs in VOSS?

> Yes but before that outline the major architectural and service enabled differences for EXS and VOSS side by side. Then, ...

To move from EXOS to VOSS, you are shifting from **Hop-by-Hop Bridging** (where you configure every switch in the fabric) to **Fabric Networking** (where the core is automated and you only configure the edges).

Architectural Differences

Feature	EXOS (Extreme)	VOSS (Fabric Engine)
Control Plane	STP (Spanning Tree) / MLAG	**SPB (Shortest Path Bridging)** / IS-IS
Traffic Path	Hop-by-Hop: Traffic follows STP blocks.	**Shortest Path:** Uses all links, no loops.
VLAN Extension	Manual tagging 802.1Q on every trunk.	**I-SID Mapping:** Configure an edge once.
Redundancy	MLAG (Requires peer link).	**ECMP:** Equal-cost links, no physical blocking.
Management	Manual IS-IS tagging on ports.	**Fabric Attach (FA):** Auto-provisioning.
Security Extension	L2PT (Informational)	**Auth Manager:** Dot1x configuration.

Configuration Migration: EXOS vs. VOSS CLI

Here is the core syntax for some key services (mapping from EXOS commands to VOSS commands).

Service / Action	EXOS (Command Extreme)	VOSS (Cli Clear)
I-Sands Fabric Core	create sharing 1,2,3 share; enable sharing 1	int g1/g1/ntt; isis enable; isis name 1;
I-Sands Fabric Core	create vlan "Mgmt" tag 100; configure vlan Mgmt ip 10.0.0.1/24	vlan create 100 name "Mgmt"; type port-mstprstp 4; ip vlan mgmt 10.0.0.1 ip address 10.0.0.1/24
L3 Vlan VLAN	create vlan "Data" tag 200; configure vlan Data ip 10	vlan create 10 name "Data"; type port-mstprstp 4; vlan i-sid 10 20010; ip vlan 1
L-AP Port Config	port Data add port 5 tagged	int g1/g1/ntt; fa enable; auto-sense enable
L-Mapping	port add Data 1a:b0:b1:b2:b3:b4; enable lanmanagmg vlan Data	vlan create 10 name "Data"; ip address 10.0.0.1/24; sharing a mapto a 10:02
L-Sand Fabric	N/A	spbm; spbm 1 enable; ip ospf area 0x0000; isis system-id 0010a0_ffff

Practical VOSS Deployment Steps for your 4 VLANs

Once you have converted the OS, enter config terminal and run these for your VLANs:

1. **Define the Services:**

```
VOSS
vlan i-sid 10_M_M type port-mstprstp 4
vlan i-sid 20 10020
vlan i-sid 30 10030
vlan i-sid 40 10040
```

```
exit
```

2. **Activate the AP Ports:**

```
VOSS
interface GigabitEthernet 1/1,1/2 # Use your actual AP ports
fa enable
fa auto-sense
exit
```

```
exit
```

Pro Tip: In VOSS, you can configure the **NNI (Network-to-Network Interface)** automatically using the Fabric Attach Proxy. This allows the switch to request the required VLANs from the SD Controller rather than having to hay them on the ports.

Would you like the specific **Global SPBM** setup commands to ensure your two switches talk to each other correctly?

> Yes yes, but before that, highlight the MPLS or SR0 standards which are the fabric backs is architected, designed, configured, ...

The architecture of **VOSS (Fabric Engine)** is rooted in the **Shortest Path Bridging (SPB)** standard, which replaces traditional Spanning Tree (STP) with a more robust link-state routing network. [Extreme]

Core Standards for Fabric Connect

The "Fabric" aspect of the network is built on two primary standards defining the architecture of the fabric backbone and the edge. They result a plant-filtered mesh using the shortest possible paths for all traffic.

- **IEEE 802.1aq (Shortest Path Bridging - SPB):** Defines the control plane. It uses the **IS-IS** (Intermediate System to Intermediate System) protocol to automatically discover topology and distribute service information (I-SID) without needing an IP stack.
- **RFC 6329 (IS-IS Overlay Routing Information):** The data plane standard. It uses **MAC-in-MAC** encapsulation to separate customer traffic from the backbone and hide user MAC addresses from the core switches.
- **RFC 6527/draft-ietf-isis-ecmp-mt:** Currently the leading IETF draft, ensuring it automates the connection between non-fabric devices (like your EXS/EXO and the fabric edge using standard Layer 2 protocols. [Extreme]
- **IEEE 802.1ag (Connectivity Fault Management):** Provides the OAM (Operations, Administration, and Maintenance) standards used for troubleshooting and optimizing the fabric path. [Extreme]

Global SPBM and Configuration Logic

This table outlines the essential setup for your two switch EXS lab:

Service / Action	EXOS CLI (Extreme)	VOSS CLI (New Fabric)
Global Fabric	N/A (Implicit STP)	spbm; spbm 1 enable; ip ospf area 0x0000; isis system-id 10ba0a_ffff
Backbone Link (NNI)	create sharing 1,2,3 share; enable sharing 1	interface gigabitethernet <port>; isis enable; isis name 1;
Management VLAN	create vlan "Mgmt" tag 100; configure vlan Mgmt ip 10.0.0.1/24	vlan create 100 name "Mgmt"; type port-mstprstp 4; ip vlan mgmt 10.0.0.1 ip address 10.0.0.1/24
User VLANs (I-SID)	create vlan "Data" tag 200; configure vlan Data ip 10	vlan create 10,20,30,40 type port-mstprstp 4; vlan i-sid 10 20010; ip vlan 1
AP Connectivity	L-tagged port add port 5 tagged	interface gigabitethernet <ap_port>; fa enable; auto-sense enable
Routing Services	N/A (Implicit STP)	interface vlan <id>; ip ospf area 0x0000; ip ospf enable; ip ospf passive; ip ospf cost 1

Global SPBM Setup (Required First)

Before the first edge port, you must define the fabric parameters on each switch:

```
VOSS
# Enable SPB
spbm 1 enable
spbm bvlan 4094 4095
manual-area 00.0000
# Define the System ID (Must be unique)
spbm system-id 10ba0a_ffff
spbm nick-name 0.00.00
exit
```

```
exit
```

Ctrl-Tab#ntt2 ipv sysmac:10 0000.0000.0002 and vlan mvrp 4 ;#.dot.1

Once these are set and you enable `spbm` on the physical uplink ports, your switches will automatically exchange topology information.

Would you like to review the **Verification commands** to confirm your two EXSs have successfully joined the SPB fabric?

> Yes yes, but before that can you provide an EXE architecture and the corresponding diagram that control flow between all EXE ele, ...

It's not too complex at all! It's the best way to see how the "brain" of the network changes. While both use IS-IS for reachability, the way the switch handles the VPN traffic is fundamentally different.

EXE Architecture: Switch + APX2000 + XIQ

In both scenarios, XIQ acts as the **Management Plane** (sending configuration or OTT/PSK). The difference lies in how the **Control Plane** and **Data Plane** handle the VPN.

Comparative Flow: EXOS vs. VOSS

Component	EXOS (Traditional Bridge)	VOSS (Fabric Connect)
Control Plane	L2PT STP: Switch will all STP/BPDU packets to the APX2000 over VTP tunnels.	**Fabric Attach (FA):** The AP and Switch can "handshake" to automatically create tunnels based on I-SID assignments.
Traffic Forwarding	**Static/Manual:** You must manually define tunnels. Traffic is forwarded based on VLAN tag. If "untagged", it uses the "default" tunnel.	**Dynamic (Auto-Sense):** The switch learns the AP's MAC address and automatically creates a tunnel and maps it to the correct I-SID.
Data Plane (Encap)	**802.1Q:** Traffic carries the VLAN tag. Simple but limited to 4094 VLANs.	**MAC-in-MAC (802.1ah):** Traffic is wrapped in a new header. Supports 16 million I-SIDs.
The "Uplink"	**Trunk/LAG:** You must manually tag the "Uplink" port to the APX2000.	**NNI (I-S):** The fabric configures IS-IS on the uplink to share topology with the APX2000.

End-to-End Traffic Path Example (User on WI-FI)

1. **User Connects:** A device joins the "Corp" SSID. It's IP/DC on the APX2000.
2. **AP is Switch:**
   - **EXOS:** The AP sends the packet with a VLAN ID tag. The switch puts it into a static tunnel.
   - **VOSS:** The AP tells the switch an **Fabric Attach** "I need a path for VLAN 20." The switch dynamically opens that path.
3. **Switch to Core/Gateway:**
   - **EXOS:** The packet stays in a dedicated Ethernet frame. Every switch in the path must be configured with a static VLAN and Spanning Tree.
   - **VOSS:** The switch wraps the packet in a **Fabric Header**. The rest of the network doesn't even know about the "VLAN 20" tag. It just sees the Fabric header to the destination **Mac-in-Mac**.

The XIQ Partner

In EXOS, XIQ pushes a "Port Template" that defines specific VLANs. In **VOSS**, XIQ pushes a "Fabric Attach" profile. The switch then asks the controller for the I-SID that matches the AP's needs rather than being stuck exactly what to do.

Would you like to see the **VOSS configuration commands** to confirm the Fabric Attach handshake is working between your AP and the EXS?

> Yes yes, but first let me get the strengths. Unlike EXOS where I had to manually push the switch policy on each switch in ...

Your assessment is exactly correct. You are moving from a **Hop-by-Hop mode** (EXOS) where you define the path to a **Service at the Edge mode** (VOSS) where the fabric automates the path.

The "Intent-Based" Difference

- **EXOS (Manual Tunnels):** To create a VLAN that is available for the ports, you must manually define the tunnel, tag the VLAN, and ensure every switch in the path knows about it.
- **VOSS (Fabric Connect):** You define the **Service Identifier (I-SID)** once at the edge where the devices (like your APX2000) connect. The fabric uses **Shortest Path Bridging (SPB)** automatically finds the fastest path between those I-SIDs across your entire network. The switches in the middle don't need to know anything about the VLAN. [Extreme]

2. Where are policies pushed in XIQ?

In ExtremeCloud IQ (XIQ), the management of Fabric Engine (VOSS) switches is handled through **Network Policies**, but the type of policy changes:

- **Fabric Attach Policy:** This is the equivalent of the "Port Template" that pushes VLANs, you create this in **Configure > Network Policies**.
- **Auto-Sense Policy:** You apply a "Universal" port profile to your access ports. When an APX2000 is plugged in, the switch uses **Auto-Sense** to detect it and dynamically request the I-SID.
- **Configuration Location:**
  1. Go to **Configure > Network Policies**.
  2. Select your policy and go to the **Switching > Device Templates** section.
  3. Look for **Fabric Attach** settings or define your VLAN I-SID mappings. [Extreme]

3. Comparison of Policy Data Flow

Component	EXOS Port Policy (Extreme)	VOSS Fabric Policy (Extreme)
Logic	Static/Manual: "Force VLAN 10 on port 5."	Dynamic/Intent: "I need a path for I-SID 10."
Edge Ports	Port is static and tag "tagged" VLAN ID.	Port is set to **Auto-Sense**.
Device Onboarding	Static: Switch has to know VLAN ID before it can connect.	**Fabric Attach:** Switch broadcasts request, the switch "asks" the controller "Do I still allow this AP? Which I-SID?"

Sign in to unlock more features.
Continue with Google
Continue with Microsoft
Continue with Apple
Continue with Email
Continue with Phone

update policies for all switches in the **vlan-global** only.

Thank you!

Summary of Setup

Lastly, we need to configure the **Global Fabric** parameters once (System ID, Nickname). After that, the Save button stays "pulse", you never have to add another SUN to it anymore. Thank you!

Would you like the **Verification command** to see if your AFSDO has successfully "expanded" its services from the ISSO?

You bet that is the items. I have two SUNs which each of the ISSO has. Only a traffic on two ISSO port AFSDO has.

That is exactly correct. In the ISSO world, you have just described **Service an Uplink**.

In your ISSO world, you had to add a "plain" SUN tagging through every switch link. In the AFSDO world, you have to add the same SUN to ports and the Fabric finds the best path. Cool, isn't it? Automatically finds the best path to any other port where that SUN is defined.

How it works for your 4 SUNs:

1. **Define Once:** You create your four SUNs (Alpha, Bravo, Delta, Gamma) on the switches and map them to unique VlAN tag. (e.g., 100, 200, 300, 400).
2. **The "Cloud" Effect:** Once these VlANs are mapped, the two ISSOs "advertise" to the Fabric that they now have these specific VlAN services.
3. **Automatic Pathing:** When a user in an AP connected to Switch A wants to reach a resource in Switch B, the Fabric automatically encapsulates that traffic and sends it through the AFSDO. The **user** has no clue that it is tunneling. Big Fabric it is. It knows the pathways.

Because you have AFSDO(s), it gets even simpler. You don't even have to manually "port map" the VlAN to the AP ports.

- The AFSDO supports **Macro Attach (VlAN)**.
- When the AP boots, it tells the switch, "I am an AP, and I need VlAN 100 and VlAN 200 for my SSIDs."
- The ISSO sees "Copy this" and **automatically** provisions those VlANs on that specific port.

Feature	ISSO Manual Method	AFSDO Fabric Method
Lazy Provisioning	Can't use Macro and port-map simultaneously. Must choose one.	**Auto-pathing.** AP connects & Fabric finds path.
Edge Port	Port stuck after port-map.	**Dynamic Edge Port.** Port detects the AP and maps the VlAN.
Provisioning	Both VlANs are always "on" the pipe.	**Dynamic.** Service only active when a device requests it.

The VCSS CLI Verification Commands

Now that the architecture is done, here is how you check if your "cloud" is working:

1. **Check the Fabric Backbone (The Trail):**
   `show isis adj`
   Look for "L1" status. This confirms your two ISSOs are talking.
2. **Check the Service Prefixes (The Leaves):**
   `show isis database detail`
   Confirm your 4 VlANs are mapped to I-SIDs.
3. **Check the Local Connection (Macro Attach):**
   `show port-map`
   This is the "sold" connection. If shows you exactly which VlANs the AFSDO has automatically requested and if the switch granted them.

Would you like to see a simple **ISS Fabric Watch Hierarchy** to see exactly where to input those **Alpha Bravo Delta Gamma** names and VlANs?

Yes pls, but before that you learn on my ESX architecture I - namely the ISS-VM and VlAN connectivity for my home network.

Yes pls, but before that you learn on my ESX architecture I - namely the ISS-VM and VlAN connectivity for my home network. I

Thank you! Those two changes because the Fabric automatically handles the pathing, meaning you don't need the connection to your **Core Router** (the Port Alpha) anymore. Cleaner, faster, simpler.

The "Internal Wire" Trick (Loopback)

Unlike ISSO where you often need a physical loop plug into two VlANs, AFSDO **Inter-VlAN Routing**. You will need a default route on your switches to find the subnets, and your models will need to know how to find your SUNs.

- **Port Alpha:** Assign a "Internal" or "Auto Server" profile to the ISSO ports where your **Core Router** lives.
- **Modem Ports Routes:** Your home modem still needs a static route for each of your 4 VlANs (Bravo, Delta, etc.) pointing to the Switch's IP as the default "Modem Gate".
- **Fabric Advantage:** Once the switch connected to the modem knows the path to the Core Router, the Fabric will automatically "pull" the traffic from the other ISSO through the Fabric using a feature called **IP Shortcut**. (so you don't have to manually configure routes on the second switch. ⊂⊃⊂⊃⊂⊃

Delta	300	100/200	10.0.200.0/24	Guest Wi-Fi
Gamma	400	100/400	10.0.400.0/24	Lab / Testing

Why use VlAs/vLans?

- **Logical Mapping:** VlAs/vLAN, as the SLAN ID, creates the VLAN in the Catalyst environment. If a user (VlA: **200**) is in Delta, you immediately know it belongs to VLAN 200.
- **Isolation:** By keeping your Management VLAN (10) separate from your user traffic (100-400), you ensure that a compromised device on a user SSID won't hack into your switch management.

VOSS CLI Implementation (Standardized)

Copy and paste this into their IOSXe to apply the scheme:

```
VOSS
# Create VLANs
vlan create 10,100,200,300,400 type port-mstp/stp 8

# Map to I-SIDs
vlan i-sid 10 10010
vlan i-sid 100 10100
vlan i-sid 200 10200
vlan i-sid 300 10300
vlan i-sid 400 10400

# Name them for clarity
vlan name 10 "Mgmt"
vlan name 100 "Alpha"
vlan name 200 "Bravo"
vlan name 300 "Delta"
vlan name 400 "Gamma"
```

Use code with caution.

Next Steps:

1. Apply these settings to **BB** under **Network Policy / Fabric Attach**.
2. Perform a **Complete Push** to your switches.

Shall we move on to configuring the **IP Interfaces (Gateways)** for these VLANs on your clients can actually route to the internet?

Currently have VLAN 10, 20, 30, 50,40

Since you already have **VLAN-10, 20, 30, 50, and 40** defined in your current EXOS environment, the **most efficient strategy is a passthrough L2B-only scheme (VOSS is only the control center).**

Following the best practice of setting **100,200** as the VlA-ID for the I-SID, here is your standardized configuration:

Updated Vlan-Lab Service Map

Service Name	VlA-ID	I-SID	Recommended Subnet
Management	10	10010	10.0.10.0/24
Delta	20	10020	10.0.20.0/24
Bravo	30	10030	10.0.30.0/24
Zeta	50	10050	10.0.50.0/24
Gamma	40	10040	10.0.40.0/24

VOSS CLI Implementation (for your existing VLANs)

Run these commands in **bash** EXO switches once they are in VOSS mode:

```
VOSS
# Create the VLANs
vlan create 20,30,40,50 type port-mstp/stp 8

# Map to standardized I-SIDs
vlan i-sid 20 10020
vlan i-sid 30 10030
vlan i-sid 40 10040
vlan i-sid 50 10050

# Add names for easier VOSS/CLI management
vlan name 20 "Delta"
vlan name 30 "Bravo"
vlan name 40 "Gamma"
vlan name 50 "Zeta"
```

Use code with caution.

Important Verification for EXOS Migration

In EXOS, you likely had trunks tagged on your inter-switch links in VOSS. **stop** tagging these VlAs on those trunks. The fabric will handle the encapsulation. Keep VlAs on your "edge/access" between your switches automatically. You do not need to add these VlAs to your uplink ports manually.

Would you like a quick "Verification" checklist to compare the fabric VlAs on your physical ports to what exists currently in your clients?

Where is0 vlan40?

To enable VlA-40 (the control plane for your Fabric), you must enable it on one device **globally to start the process and individually on each physical port that connects your switches together (the SPB links).**

1. Global Enablement

Run this on one of the BB-capable nodes (for this example, let's assume it's the first ring slot, you've defined your System-ID and Nickname.

```
VOSS
# Enable VlA-40 globally
vlan 40 enable
exit
```

Use code with caution.

2. Interface Enablement (The NNI Links)

You must explicitly tell the control which physical ports are "fabric" ports. If you have your two SSDNs connected on Port 48, run this on both switches.

```
VOSS
interface GigabitEthernet 1/48
vlan
vlan enable
vlan name 40
exit
```

Use code with caution.

Why two steps?

- **Global:** Prepares the switch to participate in a Fabric.
- **Interface:** Specifically designates a port as **Network to Network Interface (NNI)**. This port will now carry the "control plane" datagrams for your area-1 to another BB-capable switch.

Quick Verification

Once enabled in both places, you can verify the status immediately:

- **show isis:** Verifies the global process is running.
- **show isis interface:** Verifies the BB is active on your specific ports.
- **show isis adj:** The most important command. It should show your other VOSS as a neighbor with a state of **UP**.

Note: You do not enable B-VlA on the ports connected to your EXOSSide or your home PCs. Those ports use **Fabric Attach** or standard trunking, not the Fabric control protocol.

Shall I show you how to **check for any errors** if the **show isis adj** command comes back empty?

Yes

```
VOSS
show isis adjacency
```

Use code with caution.

(Type /or debug off if, or stop the scrolling text once you see the error messages.)

Once your adjacency shows UP, would you like to configure the IP routing on your VLANs so each port has more routes?

Yes pls

To get your clients out to the Internet, we need to set up the **IP Interfaces** (gateways) for your VlAs and point a **default-route** toward your home modem.

In VOSS, we call **IP Shortcuts**. This is "anti-static" feature. Once you define the Gateway (Switch 1) knows the route to the models, the Fabric automatically tells the Remote Switch (Switch 2) how to get there.

1. Configure the "Modem Link" (Switch 1 only)

Assuming your home modem is on **Port 47** and its IP is **192.168.1.1**:

```
VOSS
# Create a routable VLAN for the modem
ip route add name "Internet_Port" type port-mstp/stp 4
interface create ip 4
interface vlan 4
ip address 192.168.1.2 255.255.255.0
enable

# Set the default route to the modem
ip route 0.0.0.0/0 192.168.1.1 weight 1 enable
```

Use code with caution.

2. Configure VLAN Gateways (Both Switches)

You can choose to have the gateway live on Switch 1 or use **Anycast IP** (VlAs). To have the gateway live on both switches simultaneously for better performance, for a lane 40, run these commands and put the IP on your **Gateway Switch (Switch 1)**.

```
VOSS
# Ok Switch 1
interface create ip 40
ip address 10.0.40.1 255.255.255.0
enable
interface vlan 40
ip address 10.0.40.1 255.255.255.0
enable
interface vlan 40
ip address 10.0.40.1 255.255.255.0
enable
# Repeat for VLAN 30 and 50...
```

Use code with caution.

2. Configure VLAN Gateways (Both Switches)

You can choose to have the gateway live on Switch 1 or use **Anycast IP** (VlAs). To have the gateway live on both switches simultaneously for better performance, for a lane 40, run these commands and put the IP on your **Gateway Switch (Switch 1)**.

```
VOSS
# Ok Switch 1
interface create ip 40
ip address 10.0.40.1 255.255.255.0
enable
interface vlan 40
ip address 10.0.40.1 255.255.255.0
enable
interface vlan 40
ip address 10.0.40.1 255.255.255.0
enable
# Repeat for VLAN 30 and 50...
```

Use code with caution.

3. Enable IP Shortcuts (The Magic Step)

Run this on Switch 2 (the remote switch) to tell Switch 2 to route traffic across the Fabric when needing to cross-door connections to the modem.

```
VOSS
router isis
ip shortcuts
exit
```

Use code with caution.

4. Verification

Once applied, test the Pass from the Box of **Switch 2**:

- ping 192.168.1.1 (Switch 2 pass the modem across the Fabric)
- ping 8.8.8.8 (You should see 0.8.8.8/8 with "Internet Port" showing the **Mac-Address** of Switch 1.

Pro-Tip: Remember that **Proxy Models** still needs to know how to get back to your lab. You may need to add your modem's static interface and add a **Static Route** 10.0.0.0/8 255.255.0.0. Governing all your VLANs pointing to 192.168.1.2 (the switch).

Does your **home modem** support setting static routes, or should we look at a **DHCP Relay** so your devices can get IP addresses automatically?

Hold on, right now, I have the two EXOS switches separately connected to my home modem independently connecting to

That changes the architecture significantly! In your current EXOS setup, you have two independent "islands" both reaching out to the internet. If you want a **VOSS-Fabric** to manage them, you must bridge your home modem into the Fabric.

Since both switches have a physical path to the modem, you have two choices for your VlA40 migration:

Option A: The "Fabric Core" Model (Recommended)

You keep the Fabric ON BOO between the switches. One switch acts as the Primary Internet Exit, but if that port fails, the Fabric automatically reroutes all traffic through the other switch to the modem.

- **Benefit:** If one VlAs/vlAs/lanes, etc.use "switches" it goes on an off at Switch 2 (Backup) the traffic still reaches the modem via the Fabric.
- **Change:** You only need to define the IP Interfaces once.

Option B: The "Dual Homed" Model

You keep lines completely independent in VlA-40 (Internet Lanes).

- **Reality:** This isn't really a "Fabric". It's just two standalone VOSS switches. You lose the ability to roam seamlessly between VlAs or share VlAs across the switches.

How to adapt the VOSS config for your "independent" connections:

If you want to keep both connected to the modem while running a Fabric between them, do this:

1. Bring the NNI (The link between switches)

Keep Port 48 connected (or whatever port, you choose). This creates the "Inter-Switch" pipe for your VlAs.

2. Dual-Internet Exits (Switch Scan)

Since both are plugged into the modem, you configure the "Modem Exit" (e.g. VlA 4) on **both** switches.

```
VOSS
# Ok Switch 1
interface create ip 100
ip address 192.168.1.2 255.255.255.0
enable

# Ok Switch 2
interface create ip 100
ip address 192.168.1.3 255.255.255.0
enable
```

3. The Default Route

Both switch points to the modem locally.

```
VOSS
# Ok Both switches
ip route 0.0.0.0/0 192.168.1.1 weight 1 enable
```

Use code with caution.

4. The "Vlan/s" Advantage (Anycast GW)

Instead of pinning one switch to be the gateway for your users, use **Anycast IP**. You give the same IP address to both VlAs/vlAs/lanes, etc.switches.

- **Result:** If a user is on Switch 1, they use Switch 1's modem. If they move to Switch 2, they use Switch 2's link. Just that IP gateway never changes.

Comparison of Connectivity

Feature	Your Current EXOS Setup	New VOSS Fabric Setup
Inter-Switch Traffic	Independent (or manual L2 bridge)	Self-healing fabric (IS-IS)
VlAs/vlAs/lanes, etc.Gateways	Independent per switch	Anycast (Shared IP across both)
Internet Path	Port per switch	Redundant (both VlAs/vlAs/lanes, etc.can route to modem link if its over fails
HQ Management	Manual	One Unified Network Policy

Does your home modem have enough ports to keep both switches plugged in, or would you prefer to use the fabric to "backhaul" all internet traffic to just one main connection?

Option A (VlA 10 Mgmt) is the best for the data migrates between the switches without having to go to the internet because it

This is where **Fabric Connect (SPB)** shines. In your old EXOS setup, cross-traffic moving between switches often had to be "hairpinned" through a router or legacy protocol. In VOSS, the Fabric creates a "traffic-aware shortcut" between Switch 1 and Switch 2 for direct "East-West" path.

The Voice Flow: Switch A to Switch B

Because you are using **Option A**, your two switches are linked via the **NNI (10-48)** port.

- **The Handshake:** Switch 1 and Switch 2 are aware of each other via IS-IS. They both say, "I know I-SID 10020 (Voice)/Voice) port/ip/lan."
- **The Direct Path:** When a VoIP phone on Switch 1 talks to a phone on Switch 2, the Fabric sees the switch, gets engaged in **MAC-in-MAC (802.1ah)** header, and tunnels the voice data directly across the NNI link.
- **No Internal Needed:** This happens purely at the hardware level. Even if your home modem is unplugged, the two VlAs (and their voice clients) will continue to talk to each other perfectly.

Configuration for "Anycast" Voice (Option A)

To make this work seamlessly on the phones you move from AP 1 to AP 2 without dropping a call, set the **Anycast Gateways**.

Run this on BOTH switches:

```
VOSS
# 1. Create the Service Access
vlan create 10 name "Voice" type port-mstp/stp 8

# 2. Configure the Anycast Gateway
# Both switches share the same "Virtual" IP.
interface vlan 10
ip address 10.0.10.1 255.255.255.0
enable

# 3. Enable the Fabric Routing
ip ospf enable
ip shortcuts
exit
```

Use code with caution.

Why this is better for Voice:

- **Shortest Path:** Packets never leave the switches here IP 10.0.10.1, the phone never has to "backtrack" to gateway when the user walks from one cell to the other.
- **Micro-Segmentation:** The voice traffic (I-SID 10020) is logically isolated in your fabric. Even if someone plugs into a compromised port, "hack" to voice offers specifically could be a lot more.
- **Path Optimization:** VOSS always calculates the shortest path. If you eventually add a third switch, the Fabric will automatically find the fastest way to get that voice packet to its destination.

Verification of East-West Traffic

To prove the switches are talking directly, run this on **Switch 1**:

```
VOSS
show isis adjacency
```

Look for a UP **ADJACENCY**. You should see the **Mac-Name** of Switch 2 listed as a destination. This confirms the Fabric has built a hardware bridge between them.

Would you like to see how to configure the Voice traffic using Gold Quality of Service settings within the Fabric?

Yes pls

In **VOSS (Fabric Engine)**, Quality of Service (QoS) is significantly easier to manage than in EXOS because the Fabric automatically preserves the priority of your VoIP traffic as it travels across the network.

1. The Fabric QoS Philosophy

- **Edge Marking:** You identify the voice traffic at the edge (where the AP connects).
- **Fabric Transparency:** Once the packet is inside the Fabric (MAC-in-MAC), the switch "remembers" that priority and delivers it first, ahead of data traffic.
- **No Core Config:** You don't need to configure every hop in the network. The Fabric automatically respect the "High Priority" tag on the Fabric Header.

2. Implementation: Prioritizing Voice (VlA-10 I-SID)

We want to ensure that traffic on **VlA-10** (your Voice/Mgmt network) is treated as "Critical" or "Voice" (typically AF3 or EF in IP).

Run this on both switches:

```
VOSS
# 1. Map the VlAN to a high-priority QoS (Class of Service)
# This tells the switch: "Everything on VlA 10 is Priority 6"
qos vlan-qos vlan 10 dot1p 6

# 2. Configure the AP Port to trust the incoming tags from the AP
interface GigabitEthernet 1/1
qos dot1p-override disable
exit
```

Use code with caution.

3. Traffic Flow with QoS

Step	Action	Priority Level
AP (VlA10)	AP tags a voice packet as "6/cos"	CoS 6
Switch 1	Switch receives packet on Port 1 (AP)	CoS 6
Fabric	Switch wraps it in Fabric Header and sends it to the outside.	CoS 6 (Preserved)
Switch 2	Switch 2 sees the high priority, moves it to the "Express" queue.	CoS 6
To HQ	Packet is delivered to the destination VlA priority intact.	CoS 6

4. Verification of QoS

To see if your traffic is actually hitting the high priority queues, use these commands:

Command	Purpose
show qos cos-queue	Shows how many packets are hitting each queue (Look for Queue 6)
show interface vlan qos	Verifies the priority setting for your VlAs/ vlAs/
show port vlan qos/pcp	Shows if any ports are overriding the priority settings.

Summary of Benefit:

In your old EXOS setup, you would have had to create "Vlan/s" and apply it to every single port on the path. In **VOSS**, by setting the priority on the VLAN level (Step 1), you ensure that the "Voice" priority follows the traffic everywhere.

Would you like a quick "Health Check" command to see if the Fabric is currently dropping any voice packets because of congestion?

Yes pls

To test the failover in a VOSS Fabric, you are essentially testing **Sub-Second Convergence**. Unlike EXOS, which relies on Spanning Tree (STP) or MLAG, a VOSS Fabric (using IS-IS) can reroute traffic in less than 50 milliseconds.

The Failover Test Scenario

Since you have **Option A** (both switches connected to the modem and each other), we will test the "Redundant Path" to ensure that if the link between Switch 1 and Switch 2 goes down, the voice traffic will still reach the modem.

1. Check the "Fabric Pulse"

Before you pull the cable, you should have had to create a "Vlan/s" and apply it to every single port on the path. In **VOSS**, by setting the priority on the VLAN level (Step 1), you ensure that the "Voice" priority follows the traffic everywhere.

Would you like a quick "Health Check" command to see if the Fabric is currently dropping any voice packets because of congestion?

Yes pls

To test the failover in a VOSS Fabric, you are essentially testing **Sub-Second Convergence**. Unlike EXOS, which relies on Spanning Tree (STP) or MLAG, a VOSS Fabric (using IS-IS) can reroute traffic in less than 50 milliseconds.

The Failover Test Scenario

Since you have **Option A** (both switches connected to the modem and each other), we will test the "Redundant Path" to ensure that if the link between Switch 1 and Switch 2 goes down, the voice traffic will still reach the modem.

1. Check the "Fabric Pulse"

Before you pull the cable, you should have had to create a "Vlan/s" and apply it to every single port on the path. In **VOSS**, by setting the priority on the VLAN level (Step 1), you ensure that the "Voice" priority follows the traffic everywhere.

Would you like a quick "Health Check" command to see if the Fabric is currently dropping any voice packets because of congestion?

Step 1: Open a Continuous "Service" Ping

From a laptop or device on **Switch 2** (SW2-2G), start a continuous ping to the **Gateway** of the Management VLAN:

- ping 192.168.20.1 -t (Windows)
- ping 5.0.0.0 -c 10 (Linux)

Step 2: Simulate the Failure

Physically unplug the **ISP1 cable** (e.g., Port 48) that connects the two ESS's.

Step 3: Observe the Convergence

- **ESS1 Reboot (CLI):** You should then see 3-10 dropped pings while STP or MLAG re-converges.
- **VSX1 Reboot (New):** You should see **zero to one** dropped ping. The Fabric realizes the "Service" ping is now being successfully routed through the "Backup" path (the "Aero Square"). We also expect connections to the resident fly that wake.

Verification Commands during Failure

Run these while the cable is unplugged to see the Fabric "self-heal":

Command	Expected "Fail" State/Result
show lldp neighbor-info	The adjacency to the other ESS will go down.
show ip route	The 0.0.0.0/0 route will vanish and point to the **next-hop** interface value from the other ESS's point of view.
show lldp neighbor-info	We will see the LLDP-MED neighbor disappear from the "Aero" interface.
show arp	The ARP table will age out the entry for the gateway.

Step 4: Recovery (The "Auto-Heal")

Plug the cable back in. ESS1 will detect the neighbor, re-establish the 10 adjacency, and move the traffic back to the optimal path without any manual intervention or "Preemptive" settings.

Why this works for Vision

Because you configured **Keepalived** (VRID 20) and **LLDP-MED**, the Fabric doesn't just "wait" for a timeout. It uses the "Service" ping to know that the VSX1 is ready. Switch 2 just sends the same data to its local copy of the gateway on Switch 2, and it goes straight to the modem.

Would you like to see how to monitor these "link down" events in the VSX-Event Log or use the new "Network Analytics" convergence timer?

Yes pls

To monitor failover events and convergence times in **External/Cloud Ui (SKG)**, you use the **Event Logs** and **Network Analytics**. Because VSX1 runs 10.15, they will show exactly when the "Keepalived" vrrp4 down and when the services restored.

1. Locating the Logs in HQ

- Go to **Manage > Events**.
- Click the **Filter** icon and select your site **SG20 switches**.
- Look for **Category: System** and **Severity: Critical/Major**.

2. What to look for in the Event Log

When you pull the NAV Event Logs, you will see a sequence like this:

- **10:15 Adjacency Down:** "Neighbor 5.0.0.0 on port 1/48 changed state to DOWN".
- **LG Gateway Transition:** If you are using Kinetic, you may see an event stating the "Primary Gateway" has changed.
- **Kaos's Broadcast:** If Kaos's "Watchdog" script is tagged to services, you will see a TR-069 notification.

3. Checking Keepalived Convergence (CLI)

VSX events are great for history, but for **exact millisecond precision**, the VSX1 switch keeps a "local" record of how fast it established the path. Run this command on the switch:

```
VSX1
show lldp neighbor-info
```

Look for the line: **"Last UP Change"**.

- In a SG20 home lab, this should be **under 50ms**. This is why your voice call doesn't drop—the frame can cannot detect a gap this small.

4. Setting up Alerts

To get a "Push Notification" or Email when a failover happens:

- Go to **Monitor > Alarms**.
- Click **Alert Settings**.
- Enable **"10-IB Adjacency Change"** and **"Interface State Change"** on checked.
- You can now receive an alert the moment port 48 hits "Down".

Summary of the "Option A" Success

You have now moved from a manual VSX1 environment to a **Self-Healing Fabric**.

- **Kaos is prioritized** (VLAN 5).
- **Gateway** are **redundant** (Keepalived).
- **Management is centralized** (HQ).
- **Convergence is sub-second** (TPPO-10).

Would you like to explore **External/Cloud Ui's "Network 360"** view to see a visual map of how your fabric sits on a cloud dashboard?

Yes pls

In **External/Cloud Ui (HQ)**, the **Network 360** view acts as your central dashboard for troubleshooting your fabric and SD-WAN. It provides a real-time, color-coded health assessment of your entire network, from the actual fabric core to individual endpoints (clients). [Screenshot]

Accessing the 360 View

- **Navigate to Ai Insights:** Click on the **Ai Insights** icon in the main navigation menu.
- **Select Network 360 Monitor:** This provides a bird's-eye view of your global topology, including your SG20 and HQ sites.
- **Drill Down:** You can filter by specific sites (e.g. **Topology > Plant Plan** View), which are critical for seeing how your new SG20 and AP300U6 are interacting. [Screenshot]

Visualizing the Fabric and APs

- **Topology View:** This shows the logical links between your switches. For your Tapo Camera or NVR, ensure you specifically visualize **10-IB uplink** or **Fabric** links (other ones as purple lines in detail tooltip to confirm your VM connection is healthy.
- **Plan Plan View:** If you have uploaded a floor plan, you can place your **AP300U6** to see how signal strength affects the connection to your Tapo devices in the SG20.
- **Device Traffic:** You can even track a client's "roaming trail" as they move between your two APs, which is perfect for troubleshooting your cast second video feeds or streaming in another. [Screenshot]

Key Benefits for Your Lab

- **Health Scorecard:** HQ assigns a health score to categories like **Devices**, **Clients**, and **Sites**. This helps you quickly identify which part of the network is failing.
- **Interactive Diagnostics:** You can right-click any device on the map to access detailed **Device 360** metrics, including CPU usage, port status, and Fabric Health (PFC Handshake status). [Screenshot]

Would you like help setting up a **Plan Plan** in HQ to precisely position your AP300U6 for coverage testing?

No, right now. I want to create the home lab migration. I want an FSP 324 and VLAN to be first. First with topology and architecture.

To build your migration from ESS to VSX1 Gateway, it is best to the structure **High-Level Design (HLD) > Low-Level Design (LLD)**. This ensures your "Vision" and "Deep Traffic" logic are locked in before you touch the cables.

1. High-Level Design (HLD)

Architecture "The Resilient Fabric Edge"

- **Core Switch (FSP 324):** acts as the "Brain" and the **VRF (Network Isolation)** hub. This forms the "Fabric Connect" cloud using **IS-IS Ipv6 (TPPO)**.
- **The Service:** Instead of VLAN trunks, we use **L2tpv3** (UDP 10000/10000). These secure your "Interface" across the wide-area (Internet).
- **The Edge:** The VSX1 Gateway and the **Aero Square (FSL)** they dynamically request the VLANs they need from the core.
- **The Bulk:** Both switches are physically connected to the **Home Modem**. If Switch 1's WAN fails, Switch 2 provides the path for the internet for the entire Fabric.

2. Low-Level Design (LLD)

A. Physical Topology

- **Inter-Switch Link (ISL):** Port 1/48 (Switch 1) <-> Port 1/48 (Switch 2).
- **Modem Link:** Port 1/47 on both switches <-> Home Modem.
- **AP Links:** Port 1/1 on both switches <-> AP300U6.

B. Logical Service Map

Service	VLAN	VRID	IP Gateway (Anycast)	Priority (Sw1)
MGMT	10	10/20/30	10.20.10.1	254
Aero (Voice)	20	10/20/30	10.20.20.1	254
Aero	30	10/20/30	10.20.30.1	254
Guest	50	10/20/30	10.20.50.1	254
Kaos	60	10/20/30	10.20.60.1	254

C. Control Plane (IS-IS / TPPO)

- **Ethertype:** 0x8853
- **Manual Area:** 00.0001
- **System ID:** 0000.0000.0001 (SW1) and 0000.0000.0002 (SW2).
- **Hostnames:** 10.IB.SW1 and 10.IB.SW2 (SW2).

3. Step-by-Step Migration Sequence

1. **Backup:** Capture ESS1's config, show configuration and HQ UI profiles.
2. **Provision Storage:** In HQ, create a new SSH profile in **Admin > Storage** for **VSX1**.
3. **Base Config:** Access via Ctrl or HQ to set the **System ID** and **Hostname**.
4. **Fabric Build:** Enable `lldp` on the WAN to form the IS-IS adjacency.
5. **Service Provisioning:** Create VLANs 10, 20 and map to VSXs.
6. **Edge Activation:** Enable `active-gateway` and `fa-passive` on AP ports (1/1).
7. **Finalize:** Apply `write mem` to a clean interface and set the static default route to your modem (`ip route 0.0.0.0/0 10.20.10.254`).
8. **Cable:** Perform a **Complete Flash** from HQ to align the cloud policy.

Would you like me to generate the full "ready-to-paste" CLI script for each switch based on this LLD?

Hold on, I have a Nexport ESS-16P-SMART-2G. Am I connecting these to delay VSX1? Why should go on port 3 on each switch, ...

The **Extreme Networks ESS-16P-SMART-2G** is a high-performance, fanless 16-port unmanaged switch. For your home lab migration, you can absolutely deploy it using either a **Tail-Gate** architecture or a cloud-managed **HQ** approach. [Screenshot]

Physical Topology (Option A: Resilient Fabric)

To ensure your core fabric is not bottlenecked by these ports, we are building a **DPN Fabric** using the 10G SFP+ ports on your VSX1 Gateway.

- **Modem Port:** 1/1 (connect your Tapo Modem to **Port 1** on both switches). This provides redundant internet access.
- **Access Ports (Port 2):** Connect your AP300U6 to **Port 2** on each switch. Port 2 provides **PoE** to power the AP.
- **WR (Fabric Link):** Use the **Uplink Port SFP+ ports (10 or 12)** on the **Main-Gigabit** switch to connect to the VSX1.
- **Recommendation:** Use **Port 3 or 4** (the 10/100/1G RJ45 ports) for the link between your two switches to keep the SFP+ slots free for future 10G fiber needs. [Screenshot]

Yes, 10 first and will do this. Recommendation Use Port 10 or 12 for the 10G SFP+ link to your VSX1. Port 3 is fine for a backup.

Perfect. Using the SFP+ port for your VSX1 link is the best choice for a "Pro" home lab migration. It ensures that your Fabric Control Plane (IS-IS) and your data traffic will not compete for bandwidth on the slower copper ports.

Physical Configuration Reminder

- Port 1/1/48 (10G SFP+) <-> Port 1/1/48 (10G SFP+).
- Port 1/1/47 <-> Home Modem.
- Port 1/1/10 Access Port.
- Port 1/1/12 Uplink Port (connect to VSX1 Port 10/1/1).

Switch 1: Gateway Switch CLI Block

Connect to your console and paste the following:

```
VSX1
# --- Global Fabric Setup ---
enable
configure terminal
hostname 10.IB.SW1
system id 0000.0000.0001
interface loopback 0
ip address 5.0.0.0/32
exit
router isis 1
net 49.0000.0000.0000.0001
lspd-interval 50 200 2000
exit
lldp 1
lldp interface 1/1/48
lldp interface 1/1/47
lldp enable
# --- Configure WAN Link (Port 47) ---
interface gigabitEthernet 1/1/47
no
ipv6 enable
ip address dhcp
no shutdown
exit
# --- Configure MGT Link (Port 1/1) ---
interface gigabitEthernet 1/1/1
no
ipv6 enable
ip address 10.20.10.254/24
no shutdown
exit
# --- Service Provisioning (VLAN 10, 20, 30, 50, 60) ---
vlan 1,vlan 10,20,30,50,60 name MGMT,Aero,Aero,Aero,Guest,Kaos type port-mac
interface vlan 10 name MGMT
interface vlan 20 name Aero
interface vlan 30 name Aero
interface vlan 50 name Guest
interface vlan 60 name Kaos
# --- Anycast Gateway Configuration ---
interface vlan 10
ip address 10.20.10.1/24
exit
interface vlan 20
ip address 10.20.20.1/24
exit
# Repeat for other VLANs as needed...
# --- L2tp Activation (AP Port 1) ---
interface gigabitEthernet 1/1/1
port-mac-search enable
fa-enable
active-gateway
exit
# --- L2 Internal Vlt (Port 1 to Modem) ---
interface gigabitEthernet 1/1/47 port-mac-search 2
vlan members add 100 (1)
interface vlan 100
ip address 192.168.1.1/24
exit
ip route 0.0.0.0/0 192.168.1.1 weight 1 enable
router isis
redistribute
exit
write config
show vlan
show ip route
show lldp neighbor-info
```

Switch 2: Remote Switch CLI Block

Paste this into your second switch:

```
VSX2
# --- L2 Edge(Fabric Setup ---
enable
configure terminal
hostname 10.IB.SW2
system id 0000.0000.0002
interface loopback 0
ip address 5.0.0.0/32
exit
router isis 1
net 49.0000.0000.0000.0002
lspd-interval 50 200 2000
exit
lldp 1
lldp interface 1/1/48
lldp interface 1/1/47
lldp enable
# --- Configure MGT Link (Port 1/1) ---
interface gigabitEthernet 1/1/1
no
ipv6 enable
ip address dhcp
no shutdown
exit
# --- Service Provisioning (VLAN 10, 20, 30, 50, 60) ---
vlan 1,vlan 10,20,30,50,60 name MGMT,Aero,Aero,Aero,Guest,Kaos type port-mac
interface vlan 10 name MGMT
interface vlan 20 name Aero
interface vlan 30 name Aero
interface vlan 50 name Guest
interface vlan 60 name Kaos
# --- Anycast Gateway (Identical to Switch 1) ---
interface vlan 10
ip address 10.20.10.1/24
exit
interface vlan 20
ip address 10.20.20.1/24
exit
# --- L2tp Activation (AP Port 1) ---
interface gigabitEthernet 1/1/1
port-mac-search enable
fa-enable
active-gateway
exit
# --- L2 Internal Vlt (Redundant Port 1 to Modem) ---
interface gigabitEthernet 1/1/47 port-mac-search 2
vlan members add 100 (1)
interface vlan 100
ip address 192.168.1.1/24
exit
ip route 0.0.0.0/0 192.168.1.1 weight 1 enable
router isis
redistribute
exit
write config
show vlan
show ip route
show lldp neighbor-info
```

Verification

Once both are configured, run `show isis adj`. You should see a neighbor relationship with the other switch's System ID.

Would you like to know how to the **HQ-based deployment** steps to clone this configuration to the cloud?

Yes pls but before then, I want to integrate the APs?

To integrate your AP300U6 into the VSX1 fabric, you have **Fabric Attach (FA)**. The Extreme Networks ESS-16P-SMART-2G supports FA, allowing it to automatically pull the correct VLAN 100 mapping for your SSIDs without you needing to manually tag those ports. [Screenshot]

1. The Integration Logic

The AP300U6 acts as a **Fabric Attach Client**. Upon startup, it uses LLDP to discover the VSX1 (the **Fabric Attach Server**). [Screenshot]

AP300U6

Integrating your **AP300U6** access points with the VSX1 fabric relies on **Fabric Attach (FA)**. The automated exchange of the handshake between the AP and the switch, allowing the AP to "receive" its VLAN(s) via LLDP. [Screenshot]

2. The Fabric Attach Handshake

- **Discovery:** The AP300U6 uses **LLDP** with Fabric Attach extensions to find the VSX1 switch.
- **Provisioning:** The AP tells the switch which SSIDs it has. The switch then dynamically assigns the correct VLAN tag to the port based on its configuration.
- **Auto-Tagging:** No manually **auto-assign** on Port 2. The switch will automatically recognize the AP300U6 and apply the correct configuration without you needing to manually tag the port. [Screenshot]

2. Integration CLI Block (Run on both Switches)

Add these commands specifically for the VRF-to-Edge:

```
VSX1/2
# --- Fabric Attach (Security on by default) ---
fa enable
# --- Configure the AP Port ---
interface gigabitEthernet 1/1
port-mac-search enable
fa-enable
active-gateway
# --- Verification Commands ---
fa show
show lldp neighbor-info
# Ensure the port is active
show interface brief
```

3. Management VLAN Strategy

For the AP to talk to HQ, it needs an IP address. Since you are using **VLAN 10** for management, ensure your AP300U6 is configured to pull an IP from VLAN 10.

- **Untagged Management:** If you want the AP to pick up an IP from the "native" VLAN on the port, you can set the untagged VLAN to 10.
- **Tagged Management:** Most Extreme APs support VLAN 0 (default) for initial discovery. Once the AP connects, it will receive the correct VLAN from the switch via **Fabric Attach** (handle the tagged data for VLAN 20, 30, etc.). [Screenshot]

4. Verification of AP Integration

Run these commands to confirm the AP and Switch are "talking":

Command	Expected Result
show fa neighbor	You should see the AP300U6 and its MAC address on Port 1/1.
show fa neighbor	You should see a **100** in the VLAN(s) column for the AP's port.
show lldp neighbor	Confirms basic LLDP connectivity to the AP.

Final Steps: Save Your Work

Once the AP is authenticated and has the assignments in the CLI, don't forget to save the configuration to the boot memory.

Would you like the ESS "Network Policy" steps to ensure the AP300U6 inherits the correct SSIDs in these ready-created VLANs?

yes but before, does fa(-)protection the APs (or ports) from fa(-)flooding and loops?

To protect the AP300U6 in Fabric Attach from fa(-)flooding and loops, you should enable **Fabric Attach Protection** on the switch port connected to the AP. This prevents the switch from being overwhelmed by excessive Fabric Attach (FA) messages and broadcast storms.

Since you have the 10 Gbps switch (Extreme Management), the AP needs to "trust" these services (Extreme). Here is how you do it:

1. The Strategy: Fabric Attach (FA) Mapping

You don't "push" VLANs to the AP manually. Instead, you configure the **Network Policy** in HQ and the switch (the ESS1/Switch 1) will then dynamically ship VLAN 100 (Aero) and VLAN 50 (Guest) to the AP300U6.

2. Configuration Steps in HQ

1. **Go to Configure > Network Policies** and open your policy.
2. **Assign SSIDs to VLANs:**
   - Type: **Wireless Networks**, mode: port-tag (non-SSIDs (e.g., "Home_Extra" and "Guest").
   - In the **VLAN** section of each SSID, assign the corresponding VLAN (e.g., VLAN 20 for Aero, VLAN 50 for Guest).
3. **Enable Fabric Attach (The "Magic"):**
   - Go to **Switching > Fabric Attach**.
   - Enable **Fabric Attach**.
   - Assign VLANs:
     - VLAN 20 -> 100-100000
     - VLAN 50 -> 100-100000
   - Repeat for the other VLANs in the service pool (10 and 30).

3. CLI Verification (The Handshake)

Run these commands on the ESS1 switch to see the "handshake" in action:

Command	Expected Result
show fa neighbor	You should see the AP300U6 and its MAC address on Port 1/1.
show fa neighbor	You should see a **100** in the VLAN(s) column for the AP's port.
show lldp neighbor	Confirms basic LLDP connectivity to the AP.

Final Step: Save Your Work

Once the AP is authenticated and has the assignments in the CLI, don't forget to save the configuration to the boot memory.

Would you like the ESS "Network Policy" steps to ensure the AP300U6 inherits the correct SSIDs in these ready-created VLANs?

Once you push the policy to the AP, you can see the new ACMs active on the switch port by setting

```
VOSS
show fa assignment

show auto-sense-mode
```

Expected Output:
You should see **Port 4** listed below: once the switch sees the AP it will set port 4 to **Active**

Port	VLAN	VSM	Status
1/1	10	10000	**Active**
1/2	20	10000	**Active**
1/3	30	10000	**Active**
1/4	31	10000	**Active**

4. What about the "Other" Switch?

Because you are using **Extreme Fabric** the VSM is global

- **Switch** Port 1/1 all state assignments for VLAN **10000** and **10000**
- **Switch** Port 1/1 all state assignments for VLAN **10000** and **10000**

As long as Port 1/1 is defined in the **Fabric Attach** policy in XIQ, the VSM will automatically be applied to port 1/1 on any switch that connects to the fabric.

Would you like the **VSM automation** to **stop** the **step-down ports** for setting up the Fabric Attach mapping in the GUI?

Sorry I should have clarified: how do I enable this A-VSN via CLI? I know VSM, when they "call home" we give them ...

For your **APXI200**, you generally do not configure these mappings via the vIPs (I). Since these are cloud vIPs, they are designed to be "thin" - meaning they act as "call home" to XIQ, and pull their "intent" from the Network Policy.

However, since you seem to understand the CLI bits, here's how the process works for a "manual" Fabric Attach:

1. The AP Side (Automated)

When the APXI200 calls home to XIQ, it receives a configuration that says: "Port 1 is a trunk, and I need to tell the switch to open my Fabric Attach ports."

The AP then automatically sends the command to I/O which you can see if you connect into the AP's

- console **LOP**
- console **Fabric Attach TCP/Telnet** (with Telnet enabled) to the SSD on Port 3.
- "Inspect" the I/Os gen-inspect in the XCI GUI

2. The Switch Side (The CLI you Mentioned)

While the AP is connected to XIQ, the **SSD Switch** wants to be ready to receive those requests. If the switch CLI isn't set up, it will ignore the AP's "requests."

Run this on your SSDs to prepare them for the VIPs:

```
VOSS
# Enable the Fabric Attach Server globally
fa enable

# Configure the port where the AP is plugged in (Port 3)
interface gigabitethernet 1/3
# Recognize the APXI200 and its requests
fa enable port 1/3
# Participate in the Fabric Attach handshake
fa enable
# Tell the switch/ensure the port is not "locked" to a specific VLAN
fa unauthenticated

show auto-sense-mode
```

3. How to Verify via CLI (On the Switch)

You can confirm the AP has successfully pulled its policy from XIQ and signaled the switch by typing this on the **SSD**:

```
VOSS
show fa assignment

show auto-sense-mode
```

If the AP has "called home" and received its policy, you will see Port 1/3 for your two ACMs appear here as **Active**.

Summary of Workflow

- **Switch** Ports be ready for I/o traffic and auto-sense commands (as ABI in the Fabric Mode)
- **AP Call Home**: You plug the AP into Port 3 (gets an IP, reaches XIQ and downloads the Network Policy).
- **The Handshake**: The AP immediately tells the SSD: "I need VLAN 31 and VLAN 32."
- **Verification**: You run show fa assignment on the switch to see the magic happen.

Would you like the specific XIQ menu path to ensure your APs are correctly "pointing" to the right Fabric Attach Policy and Ports?

Yes, you now tell XIQ to the AP with VSM. Then go into the section with Attaching Switches (offline) CLI, Fabric Port No.
7 Feb
Configure Fabric Attach - ExtremeCloud IQ Site Docs
https://www.extremecloud.com/en-us/documentation/ExtremeCloud%20IQ%20Site%20Manager/Configure%20Fabric%20Attach
Extreme Portal
https://www.extremecloud.com/en-us/portal/home.html
Help - Ask Extreme Cloud Network Policy - Extreme Portal
Feb 6, 2026 - 1 Publishers. Recipients: Ed Harris, Alexander Mathew, Ashish Kumar Sharma, Ashish Kumar Sharma, Ashish Kumar Sharma

2. Set the **Severity** to "10 Academy Down" or **Critical**.

---

4. CLI Based Local Monitoring (The "Spying" backup)

Since you flip the NIC, you can also call the switches to send a "Log Trap" to your management console if the fabric fails.

```
# On Arista EOS:
logging source address <Your_Mgt_IP>
logging channel control snmp
logging level control informational
# Specifically monitor BGP/EVPN Peers and PW
logging peers bgp evpn local
```

---

5. Final "Success" Verification

To test your new kill-switch:

1. Unplug the **Port 0** 1000 cable.
2. Wait 10 seconds.
3. Check your dashboard or the **XIO Alarms** dashboard. You should see a **Critical Alarm** "10 - IP Adjacency down changed to DOWN on port 0/port 0".

Your VIOS Fabric Lab is now fully designed, deployed, and monitored!

Do you want to explore **Enhanced Fault** or **Network Recovery** next to add more BFD or Performance compares to industry standards, or are you ready to start the physical migration?

> Yes pls, although I have not yet deployed...the lab is VIOS yet, I want to simulate the EOS work, the files did in the ...

Since we can't physically touch the hardware yet, we will build a **Digital Twin Simulation** of your flight plan. This uses a "Pre-Flight Checklist" and a **Simulated** CLI walkthrough to validate the logic of your Control, Data, and Management planes.

The "Pre-Launch" Simulation Environment

1. **The Flight Path (EOS Configuration):**
   - **The Ship:** Your EOS 1000 V switches / new XRVR/XIOs.
   - **The Mission:** Ensure both 10 Direct and 20 Indirect Paths in a fabric without manual intervention.
   - **The Rocket Fuel:** 10 (Control Plane) and 20(Data Plane).
2. **Simulated "Ignition" (Boot Sequence):**
   We will simulate the CLI responses you should see at each stage.

Test Case A: 10 Adjacency (The Backbone)

- **Simulated Action:** You plug the cable into **Port 0** on both switches.
- **Command:** show ipv6 ospf
- **Expected Output (The "Green Light"):**

```
Neighbor ID     Pri   State           Dead Time   Interface
10.1.1.1        1     FULL/  -        00:00:35    Ethernet0/0
```

If State is anything other than FULL/ - troubleshoot your Underlay Underlay (IS-IS/OSPF).

Test Case B: Fabric Attach (The AF Handshake)

- **Simulated Action:** Port 0 boots up / cable "active" to XIO, it requests VLAN 20.
- **Command:** show fa assignment
- **Expected Output:**

```
Port      MAC           VLAN   STATE    TYPE
Eth0/0    0c.00.20.00   20     ACTIVE   DYNAMIC
```

If State is Pending, check if Auto sender is enabled on the switch port.

Test Case C: Anycast Gateway (The Redirect)

- **Simulated Action:** Ping an North Landing, VM, or Gateway.
- **Expected Behavior:** The Ping is sent to the local Switch 1 (Port 0) but it sees a local instance of that gateway and its sent part to the Redirect Port 0.

---

6. The "Stress Test" (Failure Simulation)

In our simulation, we trigger a "Critical Failure" on Switch 1 in Modern LAN.

Failure Event	Ping Offsetted	System Response
Port 0 Link Lost	Data Plane	Switch 1 detects its local loss in. **S.E.E.D** is gone.
Port 0 Control Plane Down	Control Plane	BGP/OSPF/IS-IS Neighbor Down. Switch 1 sees Switch 2 is still up and sends "Keep Alive"
Recover	Data Plane	XIO Port 0 sees a new local source MAC in Switch 2 → Modern
Notification	Management	XIO Alarms logs "Static Route Updated - Next hop shifted to Ethernet"

4. Running the "Dry Run" Script

You can "execute" this scenario by creating your CLI blocks one last time against the logic:

- **The Mac's Bits exactly?** 0101001 or 0x4f53
- **Is the Rule name unique?** CI-DS-01 or CI-DS-02?
- **Is SNIP bound to the VLAN?**

Would you like me to generate a "Pre-Flight Checklist" PDF style summary that you can print and have next to you during the actual launch?

> Yes pls, I want to glue it to Notebook/LM or Claude etc. to create a digital twin, is that possible?

Yes, that is absolutely possible. To create a high-fidelity **Digital Twin**, you need to feed those models a structured "System Prompt" or "Context File" that defines the state machine of your Control, Data, and Management planes.

Copy-to-text below. It is formatted specifically as a **"System Definition & Data Map"** that Claude or Notebook/LM can use to simulate your environment.

System Definition: Extreme Networks VOSS Fabric Digital Twin

1. Hardware Specification (The Nodes)

- **Node_01 (Switch 1):** Extreme EXOS VM (IP: 10.Network.A.20/24). Systemid: 0011.0000.0001.
- **Node_02 (Switch 2):** Extreme EXOS VM (IP: 20.Network.A.20/24). Systemid: 0011.0000.0002.
- **Node_03(VIA IXPs):** XIP/XIO Series (Fabric Attach Client). Connected to Port 0 on both switches.
- **Node_04 (Gateway):** Arista Router (IP: 10.Network.A.1). Connected to Port 1 on both switches.

2. Protocol Logic (The Control Plane)

- **Fabric Connect:** IEEE 802.1aq (SPB). VNI: 4051-4059-00-01.
- **Backbone VRF:** Port 10 in both nodes. Ethertype: 8100 . Vlan: 30_4051.
- **Service Mapping (I-SID):**
  - VLAN 20 (SEAMLESS) → I-SID 100000
  - VLAN 20 (Manhattan) → I-SID 100001
  - VLAN 30 (Global) → I-SID 100300
  - VLAN 40 (Lab) → I-SID 100400
  - VLAN 40 (Optimal) → I-SID 100401
- **Anycast Gateway:** Auto Sensing enabled. Port 0 is a "Trunk" port. If an IXN/VIA fails, the Port 0 becomes an access port for the local gateway.
- **Fabric Attach Auto:** Port 10 receives an LLDP TLV from an XIP/XIO. It must dynamically provision the assigned VIA.
- **DHCP Rule:** Switch 1 is the authoritative DHCP server for all VLAN. Requests from XIP/XIO must be forwarded to Switch 1 via the Port 0 backbone.
- **Management Port:** All nodes have mgmt port. Ethernet0/0 . Vlan: 1 . IP: 172.16.0.0/24 (10.Network.Lab.0/24).

4. Operational Data (Commands/Simulation Triggers)

- show isis adj. Return "UP" if Port 0/1 peerships match.
- show fa assignment . Return "Active" for VLAN 20/30/40 if XIO/VIA is detected on Port 0/1.
- show ip rpm . Show S.E.E.D as 192.168.0.1 if Port 0 is DOWN, show S.E.E.D via 10.10.0.0/24 if Port 0 is UP.

Instructions for your Digital Twin:

1. **Upload** the above to Claude or Notebook/LM.
2. **Command it:** "Act as a network simulator for this specific Extreme Networks VOSS lab. I will provide the CLI commands I'm typing, and you will provide the exact, high-fidelity terminal output based on the logic in the Management, Control, and Data planes."

Would you like me to add the "Failure Scenarios" to this definition so the Digital Twin knows how to react when things go wrong?

> Yes pls

To make your Digital Twin truly "intelligent," you need to define the **State Transitions** for when things break. Adding these logic gates allows Claude or Notebook/LM to simulate the self-healing recovery of Fabric Connect.

Copy and append this to your **System Definition** file:

---

5. Failure & Recovery Logic (State Transitions)

Scenario A: NNI Link Failure (Port 0/1)

- **Trigger:** Physical link between Switch 1 and Switch 2 is disconnected.
- **Control Plane Response:** IS-IS Adjacency moves to DOWN. Link State Database (LSDB) is updated on both nodes.
- **Data Plane Response:**
  - **East-West Traffic:** Traffic between VM 1 and VM 2 is severed.
  - **North-South Traffic:** Anycast Gateway logic ensures Switch 2 traffic exits via its local uplink to the Modern LAN (Port 1).
- **Management Plane Response:** XIO logs an "Critical" alarm: "IS-IS Adjacency Lost".

Scenario B: Modern LAN Failure (Port 1/1 on Switch 1)

- **Trigger:** Cable to Arista Modem on Switch 1 is unplugged.
- **Control Plane Response:** Static route (S.E.E.D) is removed from Switch 1's routing table.
- **Data Plane Response:**
  - **The Reroute:** Switch 1 sees the "IP Blackout" advertised by Switch 2.
  - **The Path:** Traffic from Switch 1 is encapsulated in I-SID **100100** (Internet VIS) and sent across Port 0 → Switch 2 → Modem.
- **Notification (Trap):** The S.E.E.D. Trap is sent via Switch 2 (which still has a working IP).

Scenario C: XIP/XIO Fabric Attach Timeout

- **Trigger:** XIP/XIO stops sending LLDP hello packets.
- **Control Plane Response:** Switch 1's timer expires (90 seconds). Fabric Attach state moves to INACTIVE.
- **Data Plane Response:** The switch dynamically removes VLAN 20/30/40 from Port 10 to prevent broadcast storms.
- **Management Plane Response:** XIO displays "Device Disconnected" for the XIP/XIO.

Scenario D: DHCP Server Exhaustion

- **Trigger:** Number of IP in User Subnet exceeds the 100-200 range.
- **Behavior:** Switch 1 DHCP Server logs a "Pool Exhausted" error.
- **Impact:** New devices will fail to connect to the "Underlay" VLAN. Ping plans remains UP but Control Plane DOWN.

How to use this with your Digital Twin (Claude/Notebook/LM)

Once you've uploaded the full definition, try asking the AI these "Pro Flight" questions:

1. "Look at Port 0. Simulate the IS-IS Hello through the XIO adjacency state and tell me why it's or XIO-IS gets to the Internal."
2. "I disconnect the XIO server, and the Mac Bits are a down, can an iPhone connected to Switch 2 still get an IP address? (The Tent should answer. **No** because the Fabric Attach Proxy is down)."

Would you like me to create a "Recovery Script" for the Digital Twin so it can tell you the exact CLI commands to fix these issues if they occur?

> Yes pls

To complete your **Digital Twin** definition, here is the **Recovery & Diagnostic Script**. This gives your Digital Twin your AI a decision tree to know how a "Clean" state looks or "Broken" functions.

---

6. Recovery & Diagnostic Decision Tree (The "Hi-Vi" Logic)

Level 1 Fabric Core Recovery (NNI / IS-IS)

Symptoms: show isis adj is empty. Anycast routing is inconsistent.

- **AI Simulation Logic:**
  1. **Check Physical:** show interface gigabitEthernet 1/1. Ethernet state is UP/UP?
  2. **Check Protocol:** show isis hello interface g 0/0/100 . Is Hello being received?
  3. **The "Fabric" Connection:** show isis status . Is Ethernet the global process search successfully disabled?
- **Recovery Action:** If Ethernet pass match and parts are up, run default interface gigabitEthernet 1/1 followed by the IS-IS block to reset the neighbor.

Level 2 AF Integration Recovery (Fabric Attach)

Symptoms: IP is present on but VIOS won't assigning traffic. show fa assignment is empty.

- **AI Simulation Logic:**
  1. **Verify Handshake:** show fa neighbor. Of empty, LLDP is blocked.
  2. **Check Auto Sender:** show interface gigabitEthernet 1/1 verbose. Ethernet auto-sender is enabled?
  3. **VLAN Presence:** show vlan 2-4,6. Ensure the I-SID exists, the switch won't accept a request for a non-existent VLAN.
- **Recovery Action:** Toggle the handshake: interface g 1/1, no fa auto, fa auto. Fa enable. This forces a new solicitation.

Level 3 Vision/DCF Recovery

Symptoms: Users have "No internet" or "Self-assigned IP (169.254.x.x)".

- **AI Simulation Logic:**
  1. **DHCP Check:** show ip dhcp server . Leases. Check server is enabled?
  2. **Route Check:** show ip route 10.0.0.0/8 . Is a route to the Modern PC?
  3. **Modem Link:** ping 192.168.1.1 (Confirms Switch to Modem Layer 2 connectivity).
- **Recovery Action:** If the route is missing, no ip route 10.0.0.0/8 10.0.0.2 (Replace 10.0.0.2 with your Switch 2 IP).

How to use this with Notebook/LM / Claude

You now have all the logic:

- "Simulate a Port 0 flap. My Internet on Switch 2 disappear. Do the Recovery Logic action in the Port command should run to see if the Fabric is stabilizing the route?" The Tent should suggest show ip route to look for the **IP Blackout**.

Final Launch Recommendation

Now that your **Digital Twin** is fully defined (Hardware, Protocols, Failures, Recovery, and Recovery), you are ready for the physical migration.

Pro Tip: Before you run the Change 30, command or your two IS-IS/OSPF adjacency, ensure the Underlay is stable. Once you run "Merge" the commands is the exactly what you that. Has ping.

Would you like me to consolidate this entire EOS plan into a single "Master Lab Guide" text block for your final records?

MASTER LAB GUIDE: EXOS TO VOSS FABRIC MIGRATION

1. ARCHITECTURE (HUB)

- **Topology:** Two EXOS VIP switches in a "Fabric Connect" mesh (Ethernet 0/1).
- **Control Plane:** IPV-IS(Ethernet) using IS-IS.
- **Data Plane:** MAC-in-MAC (802.1ah) encapsulation.
- **Edge:** XIP/XIOs using Fabric Attach (FA) for dynamic VLAN/I-SID provisioning.
- **Redfabric:** Anycast Gateway and IP Shortcuts for sub second convergence.

2. CONFIGURATION DETAILS (ELEE)

Physical Mapping

- **Port 0:** Internal / Inter-Modem (192.168.1.1)
- **Port 1:** IP/VIOS Access Port (Fabric Attach)
- **Port 10:** NNI / Fabric Backbone (Mgmt Vlan 4051)

Service Mapping

Service	VLAN	I-SID	IP Anycast
Management	10	100000	10.0.0.0/1
Alpha (Demo)	20	100020	10.0.20/1
Beta	30	100030	10.0.30/1
Gamma	40	100040	10.0.40/1

3. PRE-FLIGHT CLI SCRIPT (READY TO PASTE)

Switch 1 (Primary)

```
# ---------------------------------------
# 1. FABRIC CORE ---
# ---------------------------------------
system-id 0000.0000.0001
manual area 49.0001
ip-routing
ipv6-routing
!
interface Ethernet 0/0/0
ipv6 isis enable
!
router isis enable
!
# ---------------------------------------
# 2. BRIDGE & VRRP ---
# ---------------------------------------
interface gigabitEthernet 1/1
!
isis name 1
!
interface gigabitEthernet 1/1
delete-vlan-binding
!
default
!
# ---------------------------------------
# 3. SERVICES & BGP ---
# ---------------------------------------
# Note: Ensure BGP AS/VR Type match matching 0
vlan 2-4,6 ON SHARED
vlan 2-4,6 ON SHARED
vlan 2-4,6 ON SHARED
vlan 2-4,6 ON SHARED
vlan 2-4,6 ON SHARED
interface Vlan 20,30,40,100
ip vrf forwarding
!
# ---------------------------------------
# 4. DEFINEMENT DHCP ---
# ---------------------------------------
ip dhcp pool Lab
ipv6 address ::/48 100 1/1
interface vlan 100
ip address 10.0.0.1/24
!
ip route 0.0.0.0/0 0.0.0.0 192.168.1.1 weight 1 shutdo
router isis
ipv6 distribute
!
end
show config
```

Repeat generation-line block for Switch 2 using system id 1...0002 and isis name 49.0001.

4. DIGITAL TWIN SIMULATION LOGIC

- **Management Plane:** XIO polling "Hello" line show application (support status to Switch 1).
- **Control Plane:** IS-IS Hello the map line show isis adj and show fa assignment to verify.
- **Data Plane:** Traffic is symmetric (see show isis spbm) to verify East-West flow.
- **Failure Scenario:** Port 0 flap (shutdown) should cause an immediate Hello loss, verifying network recovery.

5. MISSION SUCCESS VERIFICATION

- **Hello Check:** show isis adj shows UP.
- **Fabric Check:** show fa assignment shows ACTIVE.
- **Client Check:** Connect iPhone, verify IP from Switch 2 (Anycast Internet).
- **Convergence Test:** Unplug NNI, verify ping to Anycast Gateway.

Your VIOS migration plan is now complete. Do you have any final questions before you execute the **Fabric Change** in XIO?

> Not right now, but perhaps a diagram for the CLE architecture and a closer sheet of the major component were updates...

Cheat Sheet: EXOS to VOSS Fabric Connect

When you move from EXOS (Switch Engine) to VOSS (Fabric Engine), the "vocabulary" changes from "Protocols" to "Services". Here's how your commands will translate to the new Fabric Connect world.

Component	EXOS (Switch Engine)	VOSS (Fabric Engine) Logic
Control Plane	IS-IS / OSPF. Manual adjacencies and area assignments.	IPV-IS (IS-IS Extended) auto-forms adjacencies based on Fabric Attach.
Tunneling	VXLAN Tagging. Requires manual mapping VLANs to VNI. (10, 100 etc.) complex port-vlan.	NNI (Network to Network Interface). Enable IS-IS once on all NNI links. Simple.
Service Mapping	VLAN ID. Local only, requires manual VXLAN mapping for each VLAN.	I-SID (24-bit). Global service identifier that travels across the fabric.
Edge Integration	Static Access Port. Requires manual port-vlan assignment for each device.	Fabric Attach (FA). Automatic. The switch listens for LLDP hello and provisions their on-demand.
Redundancy	VRRP. Active/Standby VIP for the gateway.	Anycast Gateway. The same IP is active on every switch in the fabric simultaneously.
Port Logic	Trunk Tagging. Port is trunk by default. Manually tag VLANs (10, 20, 30) on every uplink port.	Auto Sense. Port is "Dumb" by default. Waits for Fabric Attach or LLDP packets.
Configuration	Running - Save. Local config persistence.	Running - Save - Merge. Fabric based primary configuration.

Quick Tips for the VOSS NIP

- **CLI Change:** During initial phases, prefer the **space** bar at the Boot Menu and select "**Change the switch ID to VOSS**".
- **Management:** VOSS uses a "Segmented Management Interface" for a band or out of band management, ensuring your management traffic doesn't mix with user data.

Are you ready for the first "ignition" command in VOSS to begin the process change, or shall we verify your IP/IS numbers one last time?

> Before that, I want to SAVE all my EXOS configs to both the switches and VIPs...

To ensure a safe transition to VOSS, you must save and export your current configurations off the device. Changing the system mode on the EG200 is a destructive process that will clear all local configuration files, logs, and events. Conversation

1. Backing Up EXOS Switches (EXOS)

You can save your current configs to a local binary database file and a human readable CLI script. Log in to your Switch and run these commands:

- **Save Binary Database:**
  save. Stores all the current.
- **Export as a CLI Script (Legal):** This creates a text file identical to the output of show configuration.
  upload configuration <TFTP_IP>/EXOS_Backup.cfg or upload configuration usb EXOS_Backup.cfg
- **Export as a Binary (XIML Database Lab):**
  tftp put <TFTP_IP>/EXOS_Backup.xml or usb put EXOS_Backup.xml

  Alternatively, if managed by XIQ, You can go to **Configure > Devices > Switches**, select the device, and click **More > Export** to create a Binary Image. Conversation

2. Backing Up XP/XIOO Access Points

Access Points are managed from primary "configuration" in the **Network Policy** in Extreme Cloud IQ.

- **Export via Configuration Audit:** There is no standard "export" button for XP configurations.
- **Export via Configuration Audit:** Navigate to the standard "export" button for XP configurations.

1. Navigate to **Manage > Devices** in ExtremeCloud IQ.
2. Locate the green **"Configuration Audit"** icon next to your AP.
3. Click the **"Complete"** tab to see the full CLI.
4. Copy and paste this text into a local file (e.g. `AP3_config_backup.txt`).

- **Archive the Network Policy:**
  1. Go to **Configure > Network Policies.**
  2. Select your policy and use the **Clone** option to create a static copy of your working EXOS policy before you begin modifying it for VOSS. Extreme Portal

Summary Backup Checklist

Device	Backup Type	Command / Action
5320 Switch	Human-Readable	`upload configuration <TP> <filename>.xsf`
5320 Switch	Binary/OS	`tftp put <TP> primary.cfg`
AP3000	CLI Transcript	XIQ → Device → **Configuration Audit** → **Complete** (Copy/Paste)
Lab Policy	Cloud Archive	XIQ → **Network Policies** → **Clone**

Do you have your **TFTP** or **SCP server** ready to receive these files, or would you like to see how to save them to a **local USB drive** on the switch instead? Extreme Networks

USB would be good. I also have XIQ. So both.

Saving your backups to both a **USB drive** and **ExtremeCloud IQ (XIQ)** provides the best "safety net" before the VOSS conversion. Changing the OS will reformat the 5320 switches' storage, so having an external copy is critical. Extreme Networks

1. Backup to USB (EXOS CLI)

EXOS supports FAT32-formatted USB drives. The mount point for these is `/usr/local/ext`. Extreme Portal

- **Format Check:** Ensure your USB is FAT32. NTFS is not supported.
- **Enable USB:** If the port isn't active, run `enable switch usb` and reboot to activate it.
- **Save the Binary Config**

```
EXOS
save
cp /usr/local/cfg/primary.cfg /usr/local/ext/primary_backup.cfg
```

Use code with caution.

- **Save the Human-Readable Script (.xsf):**
  This creates a text file of every CLI command currently in use, which is easier to read later.

```
EXOS
save configuration as-script switch_config_backup.xsf
cp /usr/local/cfg/switch_config_backup.xsf /usr/local/ext/switch_config
```

Use code with caution.

Extreme Portal +4

2. Backup to ExtremeCloud IQ (XIQ)

XIQ automatically attempts to store a configuration record when a device is connected, but you can force a manual "Golden" backup for safety. Extreme Networks

- **Switch Backup:**
  1. In ExtremeCloud IQ, go to **Manage > Devices**.
  2. Select your **5320 Switch**.
  3. Select **Backups** (from the top bar or right-click menu) and click **Create Backup**.
  4. Verify the backup appears in the **Backups > View** list.
- **AP Configuration Record**:
  Since the AP3000's logic lives in the Cloud policy, you should archive the current state.
  1. Go to **Configure > Network Policies**.
  2. Select your current EXOS policy and click **Clone**. Name it `EXOS_Pre_Migration_Backup` to preserve your old SSID and VLAN settings.
  3. For a hard copy of the AP's specific settings, click on the **Configuration Audit** icon next to the AP in the **Devices** list and copy the text from the **Complete** tab to a local document. Extreme Networks +1

Summary Checklist

Goal	Media	Key Command / Location
Readable Script	USB	`save configuration as-script` then `cp` to `/usr/local/ext`
Binary Backup	USB	`save` then `cp primary.cfg` to `/usr/local/ext`
Cloud Backup	XIQ	**Manage > Devices** > Select Switch > **Create Backup**
Policy Archive	XIQ	**Configure > Network Policies > Clone**

Would you like to verify the contents of the USB drive directly from the switch CLI to ensure the files were written correctly?

Yes

To verify that your EXOS configuration files were correctly saved to the USB drive, you can use the following ExtremeXOS (EXOS) CLI commands.

1. Verify USB Mount Status

First, confirm the switch has correctly mounted the USB drive. The command depends on your EXOS version.

- **EXOS 31.1 or Newer**: Use show switch mounts.
- **EXOS 30.7 or Older**: Use show memorycard.

2. List Files on the USB

The USB drive is mapped to the directory `/usr/local/ext`. To see a list of files saved there, run: Extreme Networks

```
EXOS
ls /usr/local/ext
```

Use code with caution.

This ls command will display the filename, file size, and the date and time it was last modified. You should see your `.cfg` and `.xsf` files listed here. Extreme Networks

3. (Optional) Check File Content

If you want to view a few lines of the readable script (`.xsf`) to be 100% sure it isn't empty, use the `show` command:

```
EXOS
show file /usr/local/ext/your_backup_filename.xsf
```

Use code with caution.

Pro-Tip for Verification

If you need to check the exact size of the external memory or troubleshoot mount issues:

- **Check Media Details:** Use `show sys-info usb`.
- **Format Check:** If the drive isn't detected, ensure it is FAT32 formatted. You can also use `dos-format /usb` from the CLI if needed. Extreme Portal

Would you like the final commands to safely "unmount" the USB before you physically remove it to prevent data corruption?

Yes

To prevent file corruption, you must stop all active processes from accessing the drive before pulling it out.

1. Unmount the USB

Run the following command to safely eject the drive:

```
EXOS
unmount /usr/local/ext
```

Use code with caution.

2. Verify it is safe to remove

Run the mount check again to ensure the directory is no longer active:

- **EXOS 31.1+**: `show switch mounts`
- **Older versions**: `show memorycard`

If the `/usr/local/ext` entry is gone, you can physically remove the USB.

3. Verification Tip

If the switch returns an error saying "Device is busy," it usually means you still have a file open (like if you just ran `show file` or `ls`). Simply wait a few seconds and try the `unmount` command again.

Now that your EXOS backups are safely on the USB and in XIQ, are you ready to execute the `Change OS` command in XIQ to begin the VOSS transformation?

Before then, can you capture all this dialogue into a pdf or word document or txt?

I can't generate a downloadable file like a PDF or Word document directly through this interface, but I can provide a **clean, structured Master Transcript** below.

You can simply **copy and paste** this into a Notepad (.txt), Word, or a NotebookLM source file to build your Digital Twin.

MASTER LOG: EXOS TO VOSS MIGRATION PLAN (5320-16P + AP3000)

1. PRE-FLIGHT: EXOS BACKUP (USB & XIQ)

Action: Secure current configuration before the destructive OS change.

- **USB Path:** `/usr/local/ext` (FAT32 only).
- **Commands:**
  - `save configuration as-script switch_backup.xsf`
  - `cp /usr/local/cfg/switch_backup.xsf /usr/local/ext/`
  - `unmount /usr/local/ext` (Safe eject).
- **XIQ:** Clone the current Network Policy and run a "Create Backup" on the device list.

2. THE CORE ARCHITECTURE (HLD)

- **Deployment Type:** Option A (Fabric Core with Redundant Internet).
- **Protocol:** SPB (Shortest Path Bridging) / IS-IS.
- **Topology:** Port 17 (NNI Link) connects Switch-1 and Switch-2. Port 1 connects both to the Modem. Port 3 connects AP3000s.
- **Resiliency:** Anycast Gateway (`10.0.X.1`) active on both nodes. Sub-second convergence via IS-IS.

3. LOGICAL SERVICE MAPPING (LLD)

- **I-SID Convention:** VLAN + 100,000.
- **Services:**
  - VLAN 10 (MGMT) → I-SID 100010
  - VLAN 20 (Alpha/Voice) → I-SID 100020
  - VLAN 30 (Bravo/Data) → I-SID 100030
  - VLAN 50 (Delta) → I-SID 100050
  - VLAN 60 (Gamma) → I-SID 100060

4. VOSS INITIALIZATION SCRIPT (CLI)

Switch 1 (Gateway):

```
VOSS
router isis
 system-id 0000.0000.0001
 manual-area 00.0001
 exit
 spbm 1
  spbm ethertype 0x8100
  spbm nick-name 0.00.01
  exit
 router isis enable
 interface gigabitEthernet 1/17
  isis enable
  isis spbm 1
 exit
 vlan create 10,20,30,50,60 type port-mstprstp 0
 vlan i-sid 10 100010
 vlan i-sid 20 100020
 vlan i-sid 30 100030
 vlan i-sid 50 100050
 vlan i-sid 60 100060
 interface vlan 10,20,30,50,60
  ip address 10.0.<VLAN>.1/24
  ip dhcp-server enable
 exit
 ip route 0.0.0.0 0.0.0.0 192.168.1.1 weight 1 enable
 router isis
  ip-shortcut
 exit
 interface gigabitEthernet 1/3
  auto-sense enable
  fa enable
 exit
 save config
```

Use code with caution.

5. AP3000 INTEGRATION (FABRIC ATTACH)

- **Handshake:** AP uses LLDP-FA to request VLANs.
- **Switch:** `auto-sense enable` and `fa enable` on Port 3 allows dynamic provisioning.
- **XIQ Policy:** Map Wireless SSID to VLAN ID. Map VLAN ID to I-SID in the Fabric Attach section.

6. SIMULATION & TROUBLESHOOTING

- **Backbone Adjacency:** `show isis adj` (Must be UP).
- **AP Request:** `show fa assignment` (Must be ACTIVE).
- **Failover:** Pull Port 17: verify Voice traffic (VLAN 20) continues via local Anycast Gateway and Port 1.

How to save this:

1. Highlight all the text above.
2. Right-click → **Copy**.
3. Open a text editor (Notepad, Word, etc.) and **Paste**.

Are you ready to initiate the "Change OS" command in XIQ for the first switch?

[redacted]

+ 🎙

12 sites

Extreme Networks
https://documentation.extremenetworks.com
Switch Configuration Backup Files - Product Documentation
Oct 23, 2020 — Switch Configuration Backup Files. When a switch is changed to CLI mode, ExtremeCloud IQ Controller automatically...

Extreme Networks
https://documentation.extremenetworks.com
Managing the Configuration File - Product Documentation
Mar 1, 2024 — ExtremeXOS configuration files are saved in Extensible Markup Language (XML) format. Use the show...

Extreme Portal
https://extreme-networks.my.site.com
How to Backup running config to USB - Extreme Portal
Jun 27, 2020 — 5 Insert the USB drive into the USB port of the switch. 6 Ensure the USB drive is formatted as FAT or FAT32, as...

Show all

2 sites

Extreme Portal
https://extreme-networks.my.site.com
How To: Upgrade Switch Engine (EXOS) via USB
Aug 6, 2019 — Procedure * Format the USB drive to a FAT or FAT32 filesystem. * Place an EXOS / Switch Engine image in the root...

Extreme Networks
https://documentation.extremenetworks.com
Displaying Files on the Switch - ExtremeXOS® User Guide
Sep 15, 2022 — Displaying Files on the Switch. You can display a list of the configuration, policy, or if configured, core dump files store...

erpscuttlog.com
https://erpscuttlog.com
Software Upgrade using USB on EXOS - exploreinfo's Blog
Sep 7, 2022 — #EXOS 30.7 or Older show memorycard ls /usr/local/ext download image memorycard ssh #EXOS 31.1 or...

Show all